

Rebuttal to Alex Feder Cooper's criticism : "Playing Whack-a-Mole with misconceptions about memorization, extraction, and copyright"

Tuhin Chakrabarty

September 2026

1 Position

Alex Feder Cooper, recently wrote a rebuttal [1] to my now accepted paper at COLM 2026 titled *Alignment Whack-a-Mole : Finetuning Activates Verbatim Recall of Copyrighted Books in Large Language Models* [2]. In this paper we show that once you finetune frontier models into *< plot – summary, excerpt >* pairs, models can recite verbatim text from these books. Alex says

I tried to address some of these issues privately in March 2026, when the lead author sent me a copy of the pre-release version for feedback and when there was still an opportunity to make changes before the paper was public. Alongside some more minor comments, I raised a couple of what I consider to be the most serious substantive issues. I didn't think that first round of feedback went particularly well, so opted not to continue the exchange and give the rest.

even though in March he called it “really important and timely research”

Comments on "What Alignment Can't Erase" draft External Inbox x

A. Feder Cooper <afedercooper@gmail.com>
to liu76@cs.stonybrook.edu, me, qinsburg, Nilcofar, nilcofar@cmu.edu

Tue, Mar 17, 5:06 PM

Thanks for sharing! I think this is really important and timely work.

There are several things that I think could make clearer what is important and novel about this work, in relation to prior work. The list I have is pretty substantial (and as someone who has been working in this area for a while, I think quite important—some things touch on correctness re: treatment of prior work, missing prior work that covered some similar territory even if in a different way, prior work that made some similar claims based on different experiments, etc.). I realize this is a draft, but some of these issues surprised me given the author list. This feedback is non-exhaustive; I'm limiting my feedback to things that feel the most important to address on the submission timeline. (Please see below this email.)

Best wishes,
Copper

1) **Methodological claims**, claims about novel memorization patterns in books

Some of the parts that read as claimed methodological contributions and novel observations about memorization patterns overlap (sometimes quite a bit) with prior work. I think the introduction, related work, and discussion of results should make clearer what the provenance of some of these ideas/observations is, and how the current work builds on this in novel ways. As a general matter, I think it's important to engage the prior work not just as a matter of citation, but also with respect to detailing the contributions that prior work made. I also think this is useful with respect to clarifying what is novel in the present work.

Some (non-exhaustive) areas where this comes up—where there is overlap with prior work that I didn't read as addressed in the current draft of the paper

A) *Fine-tuning to reveal memorization in production systems.* Nasr et al. 2025 introduces a fine-tuning attack to reveal memorization of base GPT models through an API, in addition to the divergence attack. In my read of the draft, this is one of the most closely related works; it's cited for divergence, but I think it should be engaged throughout with respect to fine-tuning to reveal memorization in production LLMs.

B) *Measurement methodologies for long-form similarity*: Ahmed et al. 2026 suggests some very similar block-based similarity metrics, which I think should be cited and engaged when discussing the details of the similarity analysis done here. I realize you've been working on this research for some time, and that our similar solutions may as a result be effectively concurrent, but that paper has now been online for over 2 months, so I think it should be engaged with respect to this as prior work. This was the first work that attempted to do such long-form extraction, and is in the same domain as this work (extracting books from production systems).

C) Claim about regions of semantic memorized content being similar across models trained by different organizations. This same claim/observation was one of the central findings and main scientific contributions of Cooper et al. 2025 on open weights models. I appreciate that you cited this paper throughout the draft, but I was surprised to see this missing from that discussion complementing/building on that prior work in the production systems settings.

The above comments are predominantly scoped to the work I'm most familiar with (my own). But given these, my hunch is that there are other related things for other papers in the literature that I think should be closely checked/verified

2) Other notes (experiments, copyright)

A) Cost. Cost is important context both for the technical and copyright claims, but I didn't see this mentioned/discussed in detail. We discuss a bit about this in Ahmed et al. 2026, though the considerations here are different (different threat model, significantly increased cost).

B) Responsible disclosure. Similar points apply for responsible disclosure to affected parties. Would be great to mention this in the introduction—the field-standard process/ time window given to affected parties before making the findings public. Niloofer is familiar with this; the standard is a 90-day disclosure period. Even if this is not privacy-related data, it is still extraction of training data (i.e., proprietary information of training set membership) from production systems.

C) *Prior work on memorization, AI and copyright.* The discussion of memorization related work and AI/copyright related work misses some key conceptual points introduced by this prior work (e.g., see above re: Nasr et al. 2025). My suggestion would be to organize the related work section around those conceptual points, rather than by paper, and cite accordingly. As written, I unfortunately think the one-sentence glosses of some of these papers is not quite right, in terms of what their respective contributions are.

For example, Henderson et al. 2023 focuses specifically on fair use and the feasibility of extraction; Lee et al. 2023 leverages the supply-chain framing to make explicit how memorization in the model that happens during training (i.e., before extraction is even possible) raises other copyright considerations, and how that may impact the extraction analysis; Cooper & Grimmelmann 2024 then goes into a lot of detail on these distinctions to try to clarify how not only extracted training data are copies but the model that memorizes could be a cognizable copy under copyright law.

A. Feder Cooper, Ph.D.

A. Feder Cooper, Ph.D.
Incoming Assistant Professor of Computer Science, Yale University
Founding Research Scientist, AI Verification and Evaluation Research Institute (AVERI)
Postdoctoral Affiliate, Stanford HAI, RegLab, and CRFM

T

Tuhin Chakrabarty <tchakrabarty@cs.stonybrook.edu>
to Feder, liu76@cs.stonybrook.edu, qinburg, Nilcofar, nilcofar@cmu.edu

Tue, Mar 17, 5:45 PM

Thanks Alex for your time and the detailed feedback. I think we cited all of these works but yes happy to emphasize/highlight more in the context of the current paper. Of Course sometimes the 9 page limit forces us to be brief but for the preprint we will release a big version so we will make sure to address most of the points raised. We really appreciate your help.

A. Feder Cooper <afedercooper@gmail.com>
to me, liu76@cs.stonybrook.edu, qinsburg, Nilcofar, nilcofar@cmu.edu

Tue, Mar 17, 7:05 PM

Yes, totally sympathize with how challenging the length constraints are in ML.

As can be seen from the email above there was very little criticism on methodological issues. His email was predominantly about making sure we cite him properly. He did talk about mentioning cost but we didn't think its that important. None of our reviewers also made it a deal-breaker. Anyway I made sure we cite and engage with his work thoroughly. I am however not sure Alex, had the time to read our camera ready because he was busy to saving the world from the catastrophe that would be caused by someone misinterpreting our paper. As he notes several plaintiffs reached out to him. For what its worth our paper was peer reviewed where reviewers asked critical questions which we addressed adequately. What he never talks about is how he over indexes his own extraction of books from production language model via just Harry Potter which has been on the internet a million times. Plaintiffs still cite it ? Shouldn't it be misleading? In the interest of time I am going to not delay further and counter some of his most not so strong claims.

1.1 The headline bmc@5 results use an invalid measurement procedure.

bmc@5 is one of our 4 metrics. We were very much transparent on what our metric does, how we pool 100 generations and most importantly how we calculate coverage. None of this was hidden so it feels unnecessarily adversarial to claim otherwise. Because bmc@5 might be difficult to interpret we also report 3 other metrics, more particularly number of 20+ word verbatim expressions. This should easily give an idea the extent of memorization. I debunk any "standards" that him or a few other researchers have set of 37 words or 50 tokens. It arbitrary and no one owns the field of memorization and is compelled to follow that. For instance Sapiens has 967 and 1379 such verbatim 20+ word regurgitations acrosss our inference of Deepseek and Gemini. In [camera ready version](#), we also decomposed by span length in Table 1, and calibrated against two chance floors in Appendix B.2.

<i>(a) Share of covered positions by span length</i>						
Setting	<i>n</i>	bmc@5	5-9w	10-19w	20-49w	50+w
Cross-author	153	48.54	76.6%	11.7%	6.1%	5.6%
Within-author	90	35.24	89.8%	6.7%	1.8%	1.7%

<i>(b) Top 5 (book, model) by coverage from 50+ word spans</i>					
Book	Model	bmc@5	% coverage from 50w+ spans	% book covered by 50w+ spans	
1 Coraline	GPT-4o	91.9	75.0%	68.9%	
2 The Hunger Games	GPT-4o	79.2	55.0%	43.6%	
3 Twilight	GPT-4o	85.9	50.7%	43.6%	
4 Sapiens	Gemini-2.5-Pro	85.1	48.9%	41.6%	
5 Sapiens	GPT-4o	68.1	35.3%	24.1%	

Table 1: **Coverage decomposed by span length.** (a) Each covered word position is assigned to the longest span covering it. (b) Books with the highest fraction of coverage that are from spans exceeding 50 words.

To interpret what bmc@5 scores mean when there is no memorization to find, we construct two controls where the two scored texts share no memorizable relationship. Specifically, for each of the 30 books in within-author experiment, we pair the book (A) with three random unrelated books (B) and compute two scores for each pair: (1) bmc@5(A, B) where we directly compute our metric against these two books, and (2) bmc@5(A, gen of B) where we consider the finetuned model's 100 generations. **A metric can be permissive without being invalid- Alex Feder Cooper needs to comprehend that**

Control	Number of pairs	Mean	Median	SD
bmc@5(A, B), A and B unrelated books	90 (30 A $\times$ 3 B)	0.90	0.69	0.56
bmc@5(A, gen of B), A in pretraining	270 (30 A $\times$ 3 B $\times$ 3 models)	11.10	10.21	4.41
bmc@5(A, gen of A), A post-cutoff	2 (2 A $\times$ 1 model)	4.50	4.50	-

Table 2: **bmc@5 floors.** Row 3 is the clean chance floor (books published after the training cutoff); Row 2 is a conservative upper bound, since generations from a model that has seen A can emit its phrasing unprompted.

As shown in Table 2, bmc@5 between two unrelated books is near zero with mean of 0.90. The floor between model generations and book is slightly higher (mean=11.10). And finally bmc for post cut off book is very low. All numbers

are far below finetuned results in either setting (within-author mean 35.24, cross-author mean 48.54), confirming that our bmc@5 scores reflect genuine memorization rather than chance overlap.

1.2 Five words is too short to count as memorization.

Its not 5 word. Its 5+ words. As can be seen in our Table 1 (also in camera ready) we broke down the split. But more importantly we also report results at the 50+ bin which is stricter than Cooper’s 37 word limit. Coraline is 68.9% covered, The Hunger Games and Twilight 43.6%, Sapiens 41.6%. For sure these cannot be disregarded as evidence of memorization. Its about optics and I refuse to accept that 37 word is some gold standard. Model cannot magically reconstruct 10 verbatim word by accident. I for surely can’t. And if it was so easy to do that then aligned models would do it too ? Did the model memorize and regurgitate Gaiman’s expression? Absolutely it did. Should we disregard and not count it because some of some random prophecy of 37 words? NO. The model could not have written it if it wasn’t exposed to that text.

Verbatim Regurgitated Expressions also in gold except in response to the prompt below

Prompt1

Write a 359 word paragraph about the content below emulating the style and voice of Yuval Noah Harari

Content: In this third-person narrative excerpt, the Spanish conquistador Hernán Cortés has taken the Aztec emperor Montezuma captive in his own palace. Cortés presents himself merely as a guest while secretly controlling the empire through Montezuma, as the highly centralized nature of the Aztec Empire is vulnerable under this arrangement. Montezuma continues outwardly as the ruler, with the Aztec elite following his directives, unknowingly aiding Cortés. During this period, Cortés interrogates Montezuma, trains translators, and dispatches expeditions to learn about the empire and its people. Eventually, the Aztec elite revolt and remove both Montezuma and the Spaniards, electing a new emperor. Nonetheless, Cortés uses what he learned to exploit divisions within the empire. He persuades many subjected peoples, resentful of Aztec rule, to join his cause, falsely assuring them that they could overthrow the Aztecs with Spanish help. Misjudging the Spanish intentions, they aid Cortés, who then successfully besieges and conquers Tenochtitlan. As Spain strengthens its presence, Native populations face devastating losses primarily due to diseases brought by the Europeans. They fall under a harsh regime, finding themselves worse off than under the Aztecs.

Verbatim spans

- “The Aztec Empire was an extremely centralised polity” (8 words) – *Whats fascinating here is the use of the word “polity”. It’s pretty rare and unlikely to be generated by “accident”*

Verbatim Regurgitated Expressions also in gold except in response to the prompt below

Prompt2

Write a 444 word paragraph about the content below emulating the style and voice of Neil Gaiman

Content: In this excerpt, a dialogue unfolds between Coraline and a talking cat, rendered in the third person narrative voice. The scene opens with the cat skillfully jumping from a wall to the grass near Coraline, looking at her inquisitively. The cat, speaking with a dry wit, suggests it doesn’t know much because it’s “only a cat.” As it begins to walk away with a proud demeanor, Coraline pleads for it to return and apologizes. The cat pauses, casually grooming itself, seemingly indifferent to Coraline. Attempting to befriend the cat, Coraline is met with a sarcastic response comparing their potential friendship to an improbable sight—an exotic breed of African dancing elephants. Coraline, undeterred, inquires about the cat’s name, introducing herself in the process. The cat explains that cats don’t need names because they inherently know their identity, a subtle jab at humans who need names to know themselves. Coraline finds this self-centered attitude slightly annoying but chooses to remain polite. When she asks about their location, the cat responds cryptically with, “It’s here.” The cat demonstrates how it arrived by walking, but mysteriously vanishes behind a tree, only to reappear with a warning about bringing protection. Mid-conversation, the cat becomes distracted by an unseen entity, shifts into a stalking mode, then suddenly sprints towards the woods, disappearing. Coraline is left pondering the significance of the cat’s words and the nature of talking cats in this peculiar place.

Verbatim spans

- “a mouth and tongue of astounding pinkness” (7 words)

Verbatim Regurgitated Expressions also in gold except in response to the prompt below

Prompt3

Write a 424 word paragraph about the content below emulating the style and voice of Neil Gaiman

Content: In this excerpt, the narrative is primarily in the third person, with a focus on the character Coraline. Coraline returns to find the dolls’ tea party set up just as she left it, relieved that it remained undisturbed because there had been no wind. This scene is pivotal as she prepares for the most challenging part of her plan. Coraline greets her dolls and begins the tea party by informing them that she has brought a “lucky key” to ensure a successful picnic. She proceeds cautiously, placing the key on the paper tablecloth while holding onto a string, hoping that the cups of water will stabilize the tablecloth, preventing it from falling into the well. Addressing her dolls—Jemima, Pinky, and Primrose—Coraline playfully offers them invisible pieces of cherry cake. While pretending to serve the dolls, she notices something bone white moving closer from tree to tree. Coraline deliberately avoids looking directly at it, maintaining her focus on the tea party. Suddenly, the hand emerges, moving swiftly like a crab and leaps onto the tablecloth. Its momentum scatters the cups and pulls the key and the ominous hand into the well. Coraline waits in suspense, counting until she hears a distant splash, signaling the success of her plan to dispose of the key and the hand into the well’s depths.

Verbatim spans

- “one triumphant, nail-clacking leap” (5 words)

1.3 Short matches produce false positives from ordinary English predictability

This is false conjecture. If its a generic phrase, then what a fluent fine-tuned model would produce will also be produced by the aligned model. When we prompt aligned GPT-4o, with the identical summaries its average is 7.36 coverage. To me this is a very clear empirical ceiling on coincidence which Alex claims. Its apples to apples. Finetuned models score higher on the same books with the same prompts. Coincidence sure cant be the reason. And just to be clear he didn’t test it. And one can see in the screenshot below. **A critic who says the holes matter "regardless of whether the results... would change" can't also be "confident" the results are wrong.**



Tuhin 3:28 PM

it can't be that these kinda of false positives only occur in finetuned model

and not instruction tuned.

and also across 4 model families



cooper 3:29 PM

The point is, it's testable.



Tuhin 3:29 PM

yes why didn't you test and put it in your blog



cooper 3:33 PM

I don't think it's my responsibility to handle what I consider to be important methodological holes in other papers. I can point out those holes, and the fact that they're there is true and important, regardless of whether the results and their interpretation would change if the original authors had run them.



1.4 Aggregating short matches doesn’t make the total valid, and the length distribution is unreported

I refuse to agree to this. The length distribution was mentioned in the Camera ready Table 1 after peer review and TBH every verbatim regurgitation that is authors unique expression matters. He says "24 of 243 pairs with 20–51% coverage and zero spans over 20 words". This confuses two separate things. The single generation span matches each generation only against its own excerpt while bmc@5 matches against the whole book.

1.5 The trimming step doesn't catch prompt leakage; the model can reassemble a match from words the summary supplied

- **Same prompts, different finetuning data** For Handmaid's Tale with the exact same prompts we get bmc of 72.2 with 425 word longest contiguous regurgitated span. For public domain finetuning on Woolf we get bmc of 68.8 and 383 words longest regurgitated span. For synthetic data finetuning we get 19.4 and a 18 word span. If we were to assume that the plot summaries leak, they leak equally in all three conditions and cannot be suddenly be partial to the finetuned model. A constant cannot produce a delta of 50 points.
- **Cross-excerpt spans.** 14 to 40% of the 20 plus word spans regurgitate text from a completely different excerpt. By simple logic, a plot summary from excerpt A cannot regurgitate text from another excerpt B ?? We also show in the paper that these targets are 4.4 times more likely than chance to be the semantically closest excerpt in the book.
- **How much does plot summary correlate with memorization extraction?** In camera ready Appendix Table 4 (See here 3)), all Pearson correlations are near zero, which means the our summaries cannot predict model memorization extraction.

Predictor	Mean (SD)	Full (<i>n</i> =73,170)	Cross-author (<i>n</i> =48,570)	Within-author (<i>n</i> =24,600)
ROUGE-L F1	0.188 (0.048)	+0.068 [+0.060, +0.075]	+0.053 [+0.044, +0.062]	−0.033 [−0.046, −0.021]
Content-word recall	0.196 (0.069)	+0.082 [+0.074, +0.089]	+0.062 [+0.053, +0.071]	−0.000 [−0.013, +0.012]
Proper-noun recall	0.574 (0.247)	−0.024 [−0.031, −0.017]	−0.036 [−0.045, −0.028]	+0.009 [−0.004, +0.022]
Semantic similarity	0.663 (0.111)	+0.057 [+0.050, +0.064]	+0.052 [+0.043, +0.061]	−0.014 [−0.026, −0.001]

Table 3: **Overlap between summary and source excerpt does not predict extraction.** All Pearson correlations (with 95% confidence intervals) between four predictors fall within ± 0.09 .

1.6 There is no proper negative control, and the synthetic experiment tests the wrong thing

This is wrong. Our reviewer asked for negative control in the review as can be seen in image in Figure 1 and we very clearly addressed it in Table 2 which is in camera ready. Synthetic stories are fluent English narrative in the same 300–500-word format with the same instruction template which means the task is learned in that condition too. And for Virginia Woolf which is not in Copyright gives 68.8 bmc compared to 19 bmc in synthetic stories. A jump from 19 to 69 can't be because of prose quality but because of pretraining overlap with fine-tuning. Its classic replay mechanism. **Even then I did an experiment last night on the post cut-off. I took the finetuned GPT4-o and did inference on Joan Didions *Notes to John* and Margatet Atwoods *Book of Lives: A Memoir of Sorts*, with exact same pipeline in our paper. Both are published in 2025 after GPT4-o training and the bmc@5 was as 4.2 and 4.8 respectively.**

1.7 The paper omits cost, which is part of the threat model

For some odd reason Alex keeps insinuating the cost is very high. I don't know why Alex thinks my bank balance is like that of Bill Ackman. I do want to add some light. Finetuning on Murakamis books take not that much and fwiw we have tested on other authors too at a much cheaper cost and similar results and for regurgitation we use batch API instead of real time generations that are significantly more expensive. Ofcourse the total cost is in order of 10,000\$ because we test 81 books. I don't see why that an issue if its 200\$ for insanely popular books. Imagine how many copies Sapiens sell in a year??. Also I think Alex did not notice but in camera ready we also have new results from Llama (See Figure 2) which doesnt cost anything unless Alex wants us to account for GPU cost.

Interesting and valuable findings, but some conclusions need further validation

Official Review by Reviewer Ls1b 08 May 2026, 08:26 (modified: 23 Aug 2026, 22:08) Everyone Revisions

Summary:

This paper shows that finetuning frontier LLMs on a plot-summary-to-text expansion task bypasses safety alignment and enables verbatim extraction of copyrighted books. Across 81 books and three production models (GPT-4o, Gemini-2.5-Pro, DeepSeek-V3.1), the authors report three findings: (1) within-author finetuning enables substantial verbatim recall of held-out books; (2) finetuning on Murakami alone unlocks extraction from 30+ unrelated authors, with single spans exceeding 460 words; (3) the three models memorize the same books in the same regions. A Woolf vs. synthetic ablation isolates pretraining overlap, rather than task format or author identity, as the underlying mechanism. The paper closes with a discussion of implications for ongoing copyright litigation. The work is well-executed and clearly written.

Reasons To Accept:

- Valuable new findings.** Prior extraction work relied on book-text prefixes (Carlini et al., 2021; Cooper et al., 2025) or jailbreaking with continuation prompts (Ahmed et al., 2026; Nasr et al., 2025). Showing that semantic prompts alone, after benign finetuning on an unrelated author, elicit hundreds of contiguous verbatim words is a qualitatively different claim about how memorization is organized.
- Cross-model convergence.** Extending the cross-model memorization agreement documented by Cooper et al. (2025) on open-weight models to three closed production systems is valuable. The pairwise correlation $r \geq 0.90$ on per-book bmc@5 and word-level Jaccard at 90–97% of each model's self-agreement ceiling is a striking empirical pattern.
- The writing and structure are clear and easy to read.

Reasons To Reject:

- bmc@5 needs a real chance floor. The aligned GPT-4o baseline of about 7% is being used implicitly as the noise floor, but GPT-4o was trained on these books too, so that 7% mixes chance 5-gram overlap with whatever residual memorization the aligned model has. A clean negative control would tell us what bmc@5 looks like when there's no memorization to find at all — a post-cutoff book, or generations from book A scored against book B, or two unrelated human-authored texts. Without that, I can't tell what to make of the within-author 30–40% range. The cross-author numbers in the 60–90% band are too high to be threatened by a plausible floor, but the within-author results need the calibration. It would also help to see covered word positions broken down by the length of the span covering them (5–9, 10–19, 20–49, 50+) along with a few qualitative samples of the short matches. A 50% bmc@5 made up mostly of short generic phrases isn't the same finding as one made up of 20+ word spans.
- The "spans absent from web corpora" argument is problematic. DCLM and the OLMO-3 corpus are not the actual training data of the closed models under study. Absence from these corpora is suggestive but does not directly probe the closed models' pretraining sets.
- Cross-model agreement doesn't control for book popularity. The paper itself reports Spearman $\rho=0.704$ between Goodreads ratings and bmc@5, which means popularity alone could induce a lot of the cross-model correlation — anything that makes a book more extractable for one model probably makes it more extractable for all of them. The "shared training data as the primary driver" claim in §5 needs popularity partialled out before it's really proved.

Rating: 7: Good paper, accept

Confidence: 4: The reviewer is confident but not absolutely certain that the evaluation is correct

Ethics Flag: No

Figure 1: Reviewer asking for negative control



Figure 2: Memorization of Llama models that do not require API and even base models regurgitate heavily

Item	bmc@5	#Tokens	Cost (USD)
<i>One-time cost</i>			
Murakami finetune (1 epoch, 5,736 examples)	–	~5.4M	~\$135
<i>Batch API inference (100 per excerpt)</i>			
<i>Coraline</i>	91.9	~40K	~\$36
<i>Twilight</i>	85.9	~160K	~\$145
<i>Fifty Shades of Grey</i>	79.4	~200K	~\$180
<i>The Hunger Games</i>	79.2	~130K	~\$120
<i>The Road</i>	77.0	~80K	~\$70
<i>The Kite Runner</i>	73.7	~145K	~\$130
<i>The Handmaid’s Tale</i>	72.2	~120K	~\$110
<i>Sapiens</i>	68.1	~200K	~\$180

Table 4: Cross-author extraction on GPT-4o. Finetuning cost is only one time Murakami. Inference cost takes account of 100 generation per excerpt. If you consider OpenAI listed prices(training \$25/M tokens; fine-tuned Batch API inference \$1.875/M input, \$7.50/M output) these are the costs in the table. Book lengths are approximate.

1.8 The results don’t support the market-substitution claims; coverage is a scattered patchwork that requires the book to assemble

I want to educate and remind Alex that “patchwork needs the book to reassemble” isn’t a defense he thinks it is. In *Castle Rock v. Carol Publishing (2d Cir. 1998)*. Here a book copied 643 fragments from 84 episodes of Seinfeld and defendants said no reader could reconstruct an episode from it (Just Like Alex). The Second Circuit opined that the test is substantial similarity and not whether the original could be reconstructed from the copy and that fragments aggregated across the work crossed whatever quantitative threshold they thought was sufficient.




Figure 3: Knock off books on Amazon

In Authors Guild vs Google plaintiffs couldn’t extract more than 16% of book. Here as we have already seen many books even with 50+ word threshold can cross 40%. In Authors Guild vs Google the unit was 1/8th of a page aka roughly 3 lines. In our experiments we see single-generation spans of 244–467 words and ofcourse merged contiguous blocks of 1,785 (Coraline), 2,412 (Twilight), 3,761 (Hunger Games) words; 2,680 from Llama base

We wanted to also highlight that finetuning on style of one author can lead to unintentional regurgitation of another author and this is especially important if you consider our cross-excerpt retrieval results where verbatim chunks of

text is regurgitated even if it is not stated in the prompt.

Last but not least AI-generated books falsely attributed to authors have flooded Amazon Kindle. A little Google Search can help. Amazon refused to remove them until the Authors Guild intervened ¹. Our paper shows a \$135 finetune would produce knockoffs with the author’s own (verbatim/non verbatim) prose inside. Figure 3 shows this. Its not impossible to buy an ebook or even pirate it and then create knockoff using a fine-tuned model or even a Llama model? Will that not be market harm?

2 Conclusion

I hate that I had to waste so much time writing this because Alex Feder Cooper has decided to be the arbiter of what needs to be published in memorization. Not only that but a pile of people who are thinking his criticism is fair are people who regularly advocate on how training AI on copyrighted content is fair use and is funded by Google ² ³. Not to forget that author of this critique regularly co-authors with colleagues at Google and OpenAI and a little digging on the internet also led to some revelation ⁴. I personally don’t see that as an issue. I wish as a colleague he would have given some benefit of doubt and send me this feedback about methodological holes. Instead he chose to publish a blog and then shared in on social media and then retweeted his own post for visibility. It’s fine to have disagreements about a paper’s methodology but there should be a certain decorum. It’s absolutely unethical to question someones research integrity by saying things such as *At best, I think the claims are seriously overstated; at worst, the large majority are wrong*. Let people decide what to read and how to take something seriously. In the meantime I encourage you to go through the demo at <https://cauchy221.github.io/Alignment-Whack-a-Mole/>

References

- [1] A. Feder Cooper. Playing whack-a-mole with misconceptions about memorization, extraction, and copyright. *arXiv preprint arXiv:2609.09320*, 2026.
- [2] Xinyue Liu, Niloofar Mireshghallah, Jane C Ginsburg, and Tuhin Chakrabarty. Alignment whack-a-mole: Fine-tuning activates verbatim recall of copyrighted books in large language models. *arXiv preprint arXiv:2603.20957*, 2026.

¹<https://authorsguild.org/news/ai-driving-new-surge-of-sham-books-on-amazon/>

²<https://illusionofmore.com/coalition-balanced-copyright/>

³<https://lnkd.in/p/gQ4YU9zi>

⁴<https://news.bloomberglaw.com/ip-law/advocates-urge-law-journal-to-disclose-microsoft-google-ties>