

TUHIN CHAKRABARTY

<https://scholar.google.com/citations?user=HCmFuo8AAAAJ&hl=en>

219 W 13th Street, NY, NY 10011

<https://tuhinjbcse.github.io/>

RESEARCH INTERESTS

Human-centered AI, Natural language processing, Natural Language Generation, AI and Future of Creative Labor, Generative AI and Fair Use.

PROFESSIONAL EXPERIENCE

Stony Brook University (SUNY)

Assistant Professor (Tenure Track)

Department of Computer Science

September 2025-Present

Salesforce Research, New York

Visiting Professor, Interactive AI

May 2024 - Present

EDUCATION

Columbia University, New York

Ph.D. in Computer Science

Recipient of *Amazon PhD Fellowship for 2022,2023*

Advisor : **Smaranda Muresan**

August 2019 - 2024

GPA: 4.29

FUNDING

McGraw Hill LLC

Understanding Memorization of Training Data in Large Language Models

How AI generated books are stored in Language Models?

Jun 2026 - 2027

126,000\$

The French Publishers Association

Extracting French Copyrighted books from Base and Finetuned Language Models

Training data extraction from non-aligned models.

July 2026

24,000\$

Susman Godfrey L.L.P.

Sponsored research in lieu of In re Mosaic LLM litigation

How AI generated books dilutes the market for human authors?

Dec 2025 - 2026

70,297\$

** A previous version of my resume stated that I received an unrestricted gift for sponsored research from Susman Godfrey LLP in the amount of \$100,000. This was incorrect as the research was done for In Re Mosaic LLM litigation for which my institution was compensated in a lesser amount.

MEDIA

- *AI Has Plunged the Book Publishing Industry Into Utter Chaos* at The Wall Street Journal
- *Commentary from Atlantic CEO Nicholas Thomposon* at The Most Interesting Thing in Tech NewLetter
- *How A.I. Books Sneak Their Way Into Stores* at The New York Times
- *Is This What Comes After AI Slop?* at The Atlantic
- *The Most Famous AI Writing Tic Is Also the Most Mysterious* at The Atlantic
- *How do we spot AI slop? Lessons from the Granta controversy* at The Times of India
- *The new literary detective: Hunting AI in every sentence* at The Economic Times
- *AI accused of ‘unjust exploitation’ as bots reprint entire books* at The Telegraph
- *How AI Is Creeping Into The New York Times.* at The Atlantic
- *A.I. Is Writing Fiction. Publishers Are Unprepared.* at The New York Times
- *Who’s a Better Writer: A.I. or Humans? Take Our Quiz.* at The New York Times

- [*What if readers like AI-generated Fiction?*](#) at The New Yorker
- [*Do \(human\) readers prefer AI writers?*](#) at Marginal Revolution by Tyler Cowen
- The Hollywood Reporter [*Why AI Isn't Funny — At Least Not Yet*](#)
- The Washington Post [*Can you tell which of these break-up texts are AI-generated? Take our quiz.*](#)
- Bloomberg News [*AI Has a Way to Go to Become A Connections Puzzle Champ*](#)
- Techcrunch [*OpenAI's 'creative writing' AI evokes that annoying kid from high school fiction club*](#)
- MIT Technology Review [*What happened when 20 comedians got AI to write their routines?, Why artists are becoming less scared of AI?, AI can make you more creative—but it has limits*](#)

AWARDS, FELLOWSHIPS

- [*Outstanding Position Paper Award \(Top 1%\)*](#) at ICML 2025
- [*Best Paper Honorable Mention \(Top 5%\)*](#) at ACM CHI 2025
- NYC Media Lab x New York Times R&D [*Computational Journalism Fellowship, 2021*](#)

PUBLICATIONS(H-INDEX 26; CITATIONS 3016)

PAPERS UNDER SUBMISSION

1. **Readers Prefer Outputs of AI Trained on Copyrighted Books over Expert Human Writers**
Tuhin Chakrabarty, Jane C Ginsburg and Paramveer Dhillon
Under submission
2. **Generative AI floods and dilutes the market for books**
Tuhin Chakrabarty, Xinyue Liu, Jane C Ginsburg and Paramveer Dhillon
Under submission

REFEREED JOURNAL AND CONFERENCE PROCEEDINGS

1. **Alignment Whack-a-Mole : Finetuning Activates Verbatim Recall of Copyrighted Books in Large Language Models**
Xinyue Liu, Nilooofar Miresghallah, Jane C Ginsburg, **Tuhin Chakrabarty**
COLM 2026 ; Oral Talk at ICML 2026, MemFM workshop
2. **Measuring AI “Slop” in Text**
Chantal Shaib, **Tuhin Chakrabarty**, Diego Garcia-Olano, Byron C. Wallace
COLM 2026
3. **Death of the Novel(ty): Beyond n-Gram Novelty as a Metric for Textual Creativity**
Arkadiy Saakyan, Najoung Kim, Smaranda Muresan, **Tuhin Chakrabarty**
ICLR 2026
4. **Can Good Writing Be Generative? Expert-Level AI Writing Emerges through Fine-Tuning on High Quality Books**
Tuhin Chakrabarty and Paramveer Dhillon
ACM Conference on Human Factors in Computing Systems (CHI), 2026 (Conference)
5. **AI-Slop to AI-Polish? Aligning Language Models through Edit-Based Writing Rewards and Test-time Computation**
Tuhin Chakrabarty, Philippe Laban and Chien-Sheng Wu
COLM 2025
6. **Position: AI Safety should prioritize the Future of Work**
Sanchaita Hazra, Bodhisattwa Majumder, **Tuhin Chakrabarty**
ICML 2025 , Oral (Top 1% submission)
Outstanding Position Paper Award (Top 5%)

- 7. Can AI writing be salvaged? Mitigating Idiosyncrasies and Improving Human-AI Alignment in the Writing Process through Edits**
Tuhin Chakrabarty, Philippe Laban, Chien-Sheng Wu
ACM Conference on Human Factors in Computing Systems (CHI), 2025 (Conference)
Best Paper Honorable Mention (Top 5%)
- 8. How do Humans and Language Models Reason About Creativity? A Comparative Analysis**
Antonio Laverghetta Jr., **Tuhin Chakrabarty**, Tom Hope, Jimmy Pronchick, Krupa Bhawsar, Roger Beaty
CogSci Society 2025
- 9. Connecting the Dots: Evaluating Abstract Reasoning Capabilities of LLMs Using the New York Times Connections Word Game**
Prisha Samadarshi, Mariam Mustafa, Anushka Kulkarni, Raven Rothkopf, Smaranda Muresan and **Tuhin Chakrabarty**
Proceedings of The 2024 Conference on Empirical Methods in Natural Language Processing
- 10. Creativity Support in the Age of Large Language Models: An Empirical Study Involving Emerging Writers**
Tuhin Chakrabarty*, Vishakh Padmakumar* , Faeze Brahman, Smaranda Muresan
16th ACM Conference on Creativity & Cognition
- 11. Art or Artifice? Large Language Models and the False Promise of Creativity**
Tuhin Chakrabarty, Philippe Laban, Divyansh Agarwal, Smaranda Muresan, Chien-Sheng Wu
ACM Conference on Human Factors in Computing Systems (CHI), 2024 (Conference)
- 12. Identifying Self-Disclosures of Use, Misuse, and Addiction in Community-based Social Media Posts**
Tuhin Chakrabarty*, Chenghao Yang*, Karli R Hochstatter, Melissa N Slavin, Nabila El-Bassel, Smaranda Muresan
Findings of the 2024 Annual Conference of the North American Chapter of the Association for Computational Linguistics
- 13. Learning to Follow Object-Centric Image Editing Instructions Faithfully**
Tuhin Chakrabarty, Kanishk Singh, Arkadiy Saakyan, Smaranda Muresan
Findings of The 2023 Conference on Empirical Methods in Natural Language Processing (Long)
- 14. NORMSAGE: Multi-Lingual Multi-Cultural Norm Discovery from Conversations On-the-Fly**
Yi Fung, **Tuhin Chakrabarty**, Hao Guo, Owen Rambow, Smaranda Muresan, Heng Ji
Proceedings of The 2023 Conference on Empirical Methods in Natural Language Processing (Long)
- 15. I Spy a Metaphor: Large Language Models and Diffusion Models Co-Create Visual Metaphors**
Tuhin Chakrabarty*, Arkadiy Saakya*, Olivia Winn*, Artemis Panagopoulou, Yue Yang, Marianna Apidianaki, Smaranda Muresan
Findings of The 2023 Association of Computational Linguistics (Long)
- 16. Help me write a Poem - Instruction Tuning as a Vehicle for Collaborative Poetry Writing**
Tuhin Chakrabarty*, Visakh Padmakumar* and He He
Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)
- 17. FLUTE: Figurative Language Understanding through Textual Explanations**
Tuhin Chakrabarty, Arkadiy Saakyan, Debanjan Ghosh and Smaranda Muresan
Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)
- 18. Fine-tuned Language Models are Continual Learners**
Thomas Scialom*, **Tuhin Chakrabarty*** and Smaranda Muresan
Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)
- 19. Multitask Instruction-based Prompting for Fallacy Detection**
Tariq Alihindi, **Tuhin Chakrabarty**, Elena Musi and Smaranda Muresan
Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)
- 20. CONSISTENT: Open-Ended Question Generation From News Articles**
Tuhin Chakrabarty, Justin Lewis and Smaranda Muresan
Findings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)

21. **It's not Rocket Science: Interpreting Figurative Language in Narratives**
Tuhin Chakrabarty, Yejin Choi and Vered Shwartz
 Transactions of the Association for Computational Linguistics (TACL), 2022 (Journal)
22. **Don't Go Far Off: An Empirical Study on Neural Poetry Translation**
Tuhin Chakrabarty, Arkadiy Saakyan and Smaranda Muresan
 Proceedings of The 2021 Conference on Empirical Methods in Natural Language Processing (Long)
23. **Implicit Premise Generation with Discourse-aware Commonsense Knowledge Models**
Tuhin Chakrabarty, Aadit Trivedi and Smaranda Muresan
 Proceedings of The 2021 Conference on Empirical Methods in Natural Language Processing (Short)
24. **Metaphor Generation with Conceptual Mappings**
 Kevin Stowe, **Tuhin Chakrabarty**, Nanyun Peng, Smaranda Muresan and Iryna Gurevych
 Proceedings of The Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021) (Long)
25. **COVID-Fact: Fact Extraction and Verification of Real-World Claims concerning the COVID-19 pandemic**
 Arkadiy Saakyan, **Tuhin Chakrabarty** and Smaranda Muresan
 Proceedings of The Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021) (Long)
26. **Figurative Language in Recognizing Textual Entailment**
Tuhin Chakrabarty, Debanjan Ghosh, Adam Poliak and Smaranda Muresan
 Findings of The Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021) (Short)
27. **MERMAID: Metaphor Generation with Symbolism and Discriminative Decoding**
Tuhin Chakrabarty, Xurui Zhang, Smaranda Muresan and Nanyun Peng
 Proceedings of The 2021 Annual Conference of the North American Chapter of the Association for Computational Linguistics (Long)
28. **ENTRUST: Argument Reframing with Language Models and Entailment**
Tuhin Chakrabarty, Christopher Hidey and Smaranda Muresan
 Proceedings of The 2021 Annual Conference of the North American Chapter of the Association for Computational Linguistics (Long)
29. **Generating similes ~~effortlessly~~ like a Pro: A Style Transfer Approach for Simile Generation**
Tuhin Chakrabarty, Smaranda Muresan and Nanyun Peng
 Proceedings of The 2020 Conference on Empirical Methods in Natural Language Processing (Long)
30. **Content Planning for Neural Story Generation with Aristotelian Rescoring**
 Seraphina Goldfarb-Tarrant, **Tuhin Chakrabarty**, Ralph Weischedel and Nanyun Peng
 Proceedings of The 2020 Conference on Empirical Methods in Natural Language Processing (Long)
31. **R³: Reverse, Retrieve, and Rank for Sarcasm Generation using CommonSense Knowledge**
Tuhin Chakrabarty, Debanjan Ghosh and Smaranda Muresan and Nanyun Peng
 Proceedings of The 2020 Annual Conference of the Association for Computational Linguistics (Long)
32. **DeSePtion: Dual Sequence Prediction and Adversarial Examples for Improved Fact-Checking**
 Christopher Hidey, **Tuhin Chakrabarty**, Tariq Alhindi, Siddharth Varia, Kriste Krstovski, Mona Diab and Smaranda Muresan
 Proceedings of The 2020 Annual Conference of the Association for Computational Linguistics (Long)
33. **AMPERSAND: Argument Mining for PERSuAsive oNline Discussions**
Tuhin Chakrabarty, Christopher Hidey, Smaranda Muresan, Kathy McKeown and Alyssa Hwang
 Proceedings of The 2019 Conference on Empirical Methods in Natural Language Processing and 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Hong Kong
** Accepted as a poster at Natural Language, Dialog and Speech (NDS) Symposium, New York Academy Of Sciences (Long)*

34. IMHO Fine Tuning Improves Claim Detection

Tuhin Chakrabarty, Christopher Hidey, Kathy McKeown

Proceedings of The 2019 Annual Conference of the North American Chapter of the Association for Computational Linguistics, (NAACL) (Short)

35. Discourse Relation Prediction: Revisiting Word Pairs with Convolutional Networks

Siddharth Varia, Christopher Hidey, **Tuhin Chakrabarty**

Proceedings of The 20th Annual meeting of Special Interest Group on Discourse and Dialogue (SIGDIAL), 2019 (Long)

TEACHING

CSE 590: Human-Centered AI, Spring 2026, Stony Brook University

MENTORING

Xinyue Liu, PhD Student, Stony Brook University (Last Stop CMU LTI)

INVITED TALKS

AI, Authors, and Markets: Human-Centered Evidence for Copyright and Creative Labor, Ohio State University, February 2026

Readers Prefer Outputs of AI Trained on Copyrighted Books over Expert Human Writers, Indian Institute of Science, December 2025

Readers Prefer Outputs of AI Trained on Copyrighted Books over Expert Human Writers, University of British Columbia, July 2025

Writing in the age of AI, Narrative Understanding Workshop, NAACL May 2025

Writing in the age of AI, Cornell Information Science, October 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, University of Colorado Boulder, Feb 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, University of Michigan Ann Arbor, Feb 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, Boston University, Feb 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, Stony Brook University, March 2024

Art or Artifice? Large Language Models and the False Promise of Creativity, Stanford NLP Seminar

Art or Artifice? Large Language Models and the False Promise of Creativity, Guest Lecture at Kathy McKeown's Class COMS 6998 Language Generation and Summarization, Columbia University, September 2023

I Spy a Metaphor: Large Language Models and Diffusion Models Co-Create Visual Metaphors, UMASS NLP Group, February 2023

Controllability in Figurative Language, Guest Lecture at Kathy McKeown's Class COMS 6998 Language Generation and Summarization, Columbia University, September 2022

Fine-tuned Language Models are Continual Learners, ETH Zurich NLP group, August 2022

CONSISTENT: Controllable Open Ended Question Generation from News Articles, New York Times, January 2022

The curious case of Figurative Language University of Utah Data Science Seminar, April 2022

PROFESSIONAL ACTIVITIES

- Standing reviewer for Nature Human Behavior, PNAS.
- Area Chair for CHI 2026, COLM 2026, ACL 2026, EMNLP 2026.
- Area Chair for COLM 2025.
- Area Chair for EMNLP 2024.
- Area Chair for ACL 2024.
- Tutorial on Creative Language Generation at EMNLP 2023 with Vishakh Padmakumar, He He, and Nanyun Peng.
- Organizing committee for the 3rd Figurative Language Processing Workshop coheld at EMNLP 2022 and the 1st Shared Stories and Lessons learned workshop at EMNLP 2022
- Program Committee at ICLR 2025, ICLR 2026, UIST 2024, NAACL 2024, COLM 2024, EMNLP 2023, ACL 2023, EACL 2023, AACL 2023, ACL 2022, EMNLP 2022, AACL 2022, EACL 2023, ACL 2021, EMNLP 2021, AACL 2022, ACL 2020, EMNLP 2020 ,EMNLP 2021, NAACL 2021, *SEM 2021 , AACL 2020, *SEM 2020
- Program Committee at 1st,2nd and 3rd Workshop on Fact Extraction and Verification
- Program Committee at 4th Abusive Language Workshop.
- Program Committee at 6th Workshop on Noisy User-Generated Text (W-NUT)