

TUHIN CHAKRABARTY

<https://scholar.google.com/citations?user=HCmFuo8AAAAJ&hl=en>

219 W 13th Street, NY, NY 10011

<https://tuhinjbcse.github.io/>

RESEARCH INTERESTS

Human-centered AI, Natural language processing, Natural Language Generation, Computational Creativity.

PROFESSIONAL EXPERIENCE

Stony Brook University (SUNY)

August 2025

Assistant Professor (Tenure Track)

Department of Computer Science

Salesforce Research, New York

May 2024 - Present

Research Scientist, Interactive AI

EDUCATION

Columbia University, New York

August 2019 - 2024

Ph.D. in Computer Science

GPA: 4.29

Recipient of *Amazon PhD Fellowship for 2022,2023*

Advisor : **Smaranda Muresan**

AWARDS, FELLOWSHIPS AND MEDIA

- *Best Paper Honorable Mention (Top 5%)* at ACM CHI 2025
- NYC Media Lab x New York Times R&D *Computational Journalism Fellowship, 2021*
- The Hollywood Reporter *Why AI Isn't Funny — At Least Not Yet*
- The Washington Post *Can you tell which of these break-up texts are AI-generated? Take our quiz.*
- Bloomberg News *AI Has a Way to Go to Become A Connections Puzzle Champ*
- Techcrunch *OpenAI's 'creative writing' AI evokes that annoying kid from high school fiction club*
- MIT Technology Review *What happened when 20 comedians got AI to write their routines?, Why artists are becoming less scared of AI?, AI can make you more creative—but it has limits*

PUBLICATIONS(H-INDEX 23)

PAPERS UNDER SUBMISSION

1. AI-Slop to AI-Polish? Aligning Language Models through Edit-Based Writing Rewards and Test-time Computation

Tuhin Chakrabarty, Philippe Laban and Chien-Sheng Wu

Under Submission at COLM 2025

2. Position: AI Safety should prioritize the Future of Work

Sanchaita Hazra, Bodhisattwa Majumder, **Tuhin Chakrabarty**

Under Submission at ICML 2025

1. Can AI writing be salvaged? Mitigating Idiosyncrasies and Improving Human-AI Alignment in the Writing Process through Edits

Tuhin Chakrabarty, Philippe Laban, Chien-Sheng Wu

ACM Conference on Human Factors in Computing Systems (CHI), 2025 (Conference)

Best Paper Honorable Mention (Top 5%)

2. Connecting the Dots: Evaluating Abstract Reasoning Capabilities of LLMs Using the New York Times Connections Word Game

Prisha Samadarshi, Mariam Mustafa, Anushka Kulkarni, Raven Rothkopf, Smaranda Muresan

and **Tuhin Chakrabarty**

Proceedings of The 2024 Conference on Empirical Methods in Natural Language Processing

3. Creativity Support in the Age of Large Language Models: An Empirical Study Involving Emerging Writers

Tuhin Chakrabarty*, Vishakh Padmakumar* , Faeze Brahman, Smaranda Muresan

16th ACM Conference on Creativity & Cognition

4. Art or Artifice? Large Language Models and the False Promise of Creativity

Tuhin Chakrabarty, Philippe Laban, Divyansh Agarwal, Smaranda Muresan, Chien-Sheng Wu

ACM Conference on Human Factors in Computing Systems (CHI), 2024 (Conference)

5. Identifying Self-Disclosures of Use, Misuse, and Addiction in Community-based Social Media Posts

Tuhin Chakrabarty*, Chenghao Yang*, Karli R Hochstatter, Melissa N Slavin, Nabila El-Bassel, Smaranda Muresan

Findings of the 2024 Annual Conference of the North American Chapter of the Association for Computational Linguistics

6. Learning to Follow Object-Centric Image Editing Instructions Faithfully

Tuhin Chakrabarty, Kanishk Singh, Arkadiy Saakyan, Smaranda Muresan

Findings of The 2023 Conference on Empirical Methods in Natural Language Processing (Long)

7. NORMSAGE: Multi-Lingual Multi-Cultural Norm Discovery from Conversations On-the-Fly

Yi Fung, **Tuhin Chakrabarty**, Hao Guo, Owen Rambow, Smaranda Muresan, Heng Ji

Proceedings of The 2023 Conference on Empirical Methods in Natural Language Processing (Long)

8. I Spy a Metaphor: Large Language Models and Diffusion Models Co-Creat Visual Metaphors

Tuhin Chakrabarty*, Arkadiy Saakya*, Olivia Winn*, Artemis Panagopoulou, Yue Yang, Marianna Apidianaki, Smaranda Muresan

Findings of The 2023 Association of Computational Linguistics (Long)

9. Help me write a Poem - Instruction Tuning as a Vehicle for Collaborative Poetry Writing

Tuhin Chakrabarty*, Visakh Padmakumar* and He He

Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)

10. FLUTE: Figurative Language Understanding through Textual Explanations

Tuhin Chakrabarty, Arkadiy Saakyan, Debanjan Ghosh and Smaranda Muresan

Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)

11. Fine-tuned Language Models are Continual Learners

Thomas Scialom*, **Tuhin Chakrabarty*** and Smaranda Muresan

Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)

12. Multitask Instruction-based Prompting for Fallacy Detection

Tariq Alihindi, **Tuhin Chakrabarty**, Elena Musi and Smaranda Muresan

Proceedings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)

13. CONSISTENT: Open-Ended Question Generation From News Articles

Tuhin Chakrabarty, Justin Lewis and Smaranda Muresan

Findings of The 2022 Conference on Empirical Methods in Natural Language Processing (Long)

14. **It's not Rocket Science: Interpreting Figurative Language in Narratives**
Tuhin Chakrabarty, Yejin Choi and Vered Shwartz
 Transactions of the Association for Computational Linguistics (TACL), 2022 (Journal)
15. **Don't Go Far Off: An Empirical Study on Neural Poetry Translation**
Tuhin Chakrabarty, Arkadiy Saakyan and Smaranda Muresan
 Proceedings of The 2021 Conference on Empirical Methods in Natural Language Processing (Long)
16. **Implicit Premise Generation with Discourse-aware CommonsenseKnowledge Models**
Tuhin Chakrabarty, Aadit Trivedi and Smaranda Muresan
 Proceedings of The 2021 Conference on Empirical Methods in Natural Language Processing (Short)
17. **Metaphor Generation with Conceptual Mappings**
 Kevin Stowe, **Tuhin Chakrabarty**, Nanyun Peng, Smaranda Muresan and Iryna Gurevych
 Proceedings of The Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021) (Long)
18. **COVID-Fact: Fact Extraction and Verification of Real-World Claims concerning the COVID-19 pandemic**
 Arkadiy Saakyan, **Tuhin Chakrabarty** and Smaranda Muresan
 Proceedings of The Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021) (Long)
19. **Figurative Language in Recognizing Textual Entailment**
Tuhin Chakrabarty, Debanjan Ghosh, Adam Poliak and Smaranda Muresan
 Findings of The Joint Conference of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2021) (Short)
20. **MERMAID: Metaphor Generation with Symbolism and Discriminative Decoding**
Tuhin Chakrabarty, Xurui Zhang, Smaranda Muresan and Nanyun Peng
 Proceedings of The 2021 Annual Conference of the North American Chapter of the Association for Computational Linguistics (Long)
21. **ENTRUST: Argument Reframing with Language Models and Entailment**
Tuhin Chakrabarty, Christopher Hidey and Smaranda Muresan
 Proceedings of The 2021 Annual Conference of the North American Chapter of the Association for Computational Linguistics (Long)
22. **Generating similes ~~effortlessly~~ like a Pro: A Style Transfer Approach for Simile Generation**
Tuhin Chakrabarty, Smaranda Muresan and Nanyun Peng
 Proceedings of The 2020 Conference on Empirical Methods in Natural Language Processing (Long)
23. **Content Planning for Neural Story Generation with Aristotelian Rescoring**
 Seraphina Goldfarb-Tarrant, **Tuhin Chakrabarty**, Ralph Weischedel and Nanyun Peng
 Proceedings of The 2020 Conference on Empirical Methods in Natural Language Processing (Long)
24. **R³: Reverse, Retrieve, and Rank for Sarcasm Generation using CommonSense Knowledge**
Tuhin Chakrabarty, Debanjan Ghosh and Smaranda Muresan and Nanyun Peng
 Proceedings of The 2020 Annual Conference of the Association for Computational Linguistics (Long)
25. **DeSePtion: Dual Sequence Prediction and Adversarial Examples for Improved Fact-Checking**
 Christopher Hidey, **Tuhin Chakrabarty**, Tariq Alhindi, Siddharth Varia, Kriste Krstovski, Mona Diab and Smaranda Muresan
 Proceedings of The 2020 Annual Conference of the Association for Computational Linguistics (Long)
26. **AMPERSAND: Argument Mining for PERSuAsive oNline Discussions**
Tuhin Chakrabarty, Christopher Hidey, Smaranda Muresan, Kathy McKeown and Alyssa Hwang
 Proceedings of The 2019 Conference on Empirical Methods in Natural Language Processing and 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Hong Kong
** Accepted as a poster at Natural Language, Dialog and Speech (NDS) Symposium, New York Academy Of Sciences (Long)*

27. IMHO Fine Tuning Improves Claim Detection

Tuhin Chakrabarty, Christopher Hidey, Kathy McKeown

Proceedings of The 2019 Annual Conference of the North American Chapter of the Association for Computational Linguistics, (NAACL) (Short)

28. Discourse Relation Prediction: Revisiting Word Pairs with Convolutional Networks

Siddharth Varia, Christopher Hidey, **Tuhin Chakrabarty**

Proceedings of The 20th Annual meeting of Special Interest Group on Discourse and Dialogue (SIGDIAL), 2019 (Long)

WORKSHOP, DEMO

1. Identifying Distributional Perspective Differences from Colingual Groups

Yufei Tian, **Tuhin Chakrabarty**, Fred Morstatter, Nanyun Peng

Proceedings of The 9th International Workshop on Natural Language Processing for Social Media, NAACL 2021 (Long)

2. DiSCoL: Toward Engaging Dialogue Systems through Conversational Line Guided Response Generation

Sarik Ghazarian, Zixi Liu, **Tuhin Chakrabarty**, Xuezhe Ma, Aram Galstyan, Nanyun Peng

Proceedings of The 2021 Annual Conference of the North American Chapter of the Association for Computational Linguistics (Demo)

3. Pay “Attention” to your Context when Classifying Abusive Language

Tuhin Chakrabarty, Kilol Gupta and Smaranda Muresan

Proceedings of The 3rd Abusive Language Workshop, ACL 2019 (Long)

4. The Answer is Language Model Fine-tuning

Tuhin Chakrabarty and Smaranda Muresan

Proceedings of The 13th International Workshop on Semantic Evaluation, NAACL-HLT, Minneapolis, NAACL 2019 (Short)

5. Robust Document Retrieval and Individual Evidence Modeling for Fact Extraction and Verification

Tuhin Chakrabarty, Tariq Alihindi and Smaranda Muresan

Proceedings of The 1st Workshop on Fact Extraction and Verification, EMNLP 2018 (Short)

INVITED TALKS

Writing in the age of AI, Narrative Understanding Workshop, NAACL May 2025

Writing in the age of AI, Cornell Information Science, October 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, University of Colorado Boulder, Feb 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, University of Michigan Ann Arbor, Feb 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, Boston University, Feb 2024

Knowledge-enhanced LLMs and Human-AI Collaboration Frameworks for Creativity Support, Stony Brook University, March 2024

Art or Artifice? Large Language Models and the False Promise of Creativity, Stanford NLP Seminar

Art or Artifice? Large Language Models and the False Promise of Creativity, Guest Lecture at Kathy McKeown’s Class COMS 6998 Language Generation and Summarization, Columbia University, September 2023

I Spy a Metaphor: Large Language Models and Diffusion Models Co-Create Visual Metaphors, UMASS NLP Group, February 2023

Controllability in Figurative Language, Guest Lecture at Kathy McKeown’s Class COMS 6998 Language Generation and Summarization, Columbia University, September 2022

Fine-tuned Language Models are Continual Learners, ETH Zurich NLP group, August 2022

CONSISTENT: Controllable Open Ended Question Generation from News Articles, New York Times, January 2022

The curious case of Figurative Language University of Utah Data Science Seminar, April 2022

Controlled Generation for Figurative Language and Argument Reframing, ETS Research, November 2021

PROFESSIONAL ACTIVITIES

- Standing reviewer for Nature Human Behavior, PNAS.
- Area Chair for COLM 2024.
- Area Chair for EMNLP 2024.
- Area Chair for ACL 2024.
- Tutorial on Creative Language Generation at EMNLP 2023 with Vishakh Padmakumar, He He, and Nanyun Peng.
- Organizing committee for the 3rd Figurative Language Processing Workshop coheld at EMNLP 2022 and the 1st Shared Stories and Lessons learned workshop at EMNLP 2022
- Program Committee at UIST 2024, NAACL 2024, COLM 2024, EMNLP 2023, ACL 2023, EACL 2023, AACL 2023, ACL 2022, EMNLP 2022, AAAI 2022, EACL 2023, ACL 2021, EMNLP 2021, AAAI 2022, ACL 2020, EMNLP 2020, EMNLP 2021, NAACL 2021, *SEM 2021, AAAI 2020, *SEM 2020
- Program Committee at 1st, 2nd and 3rd Workshop on Fact Extraction and Verification
- Program Committee at 4th Abusive Language Workshop.
- Program Committee at 6th Workshop on Noisy User-Generated Text (W-NUT)

INTERNSHIP EXPERIENCE

Research Intern, Google DeepMind

August 2023- November 2023

Structured Explanations for Automatic Evaluation of Attribution with Large Language Models.

Mentors : **David Reitter and Hannah Rashkin**

Research Intern, Salesforce Research

May 2023 - August 2023

Human-centered evaluation for Creative Writing.

Mentors : **Philippe Laban and Chien-Sheng Wu**

Summer Intern PhD, Allen Institute of Artificial Intelligence

May 2021 - August 2021

Explanatory models for Narrative Understanding involving figurative language.

Mentors : **Vered Shwartz and Yejin Choi**

Summer Intern PhD, IBM Research

April 2020 - July 2020, YorkTown Heights, NY

Generative models for Question Answering.